

Customer review rating prediction

By Simantini Patil

Introduction

Many companies today measure customer satisfaction. Some of the quick ways of collecting the feedback from customers is through their reviews or by surveying them. These are very important customer service metrics, but are not typically being used to improve operations or help reduce customer churn.

Instead of waiting until the customer interaction is over for feedback/ review, companies can predict how likely an operation is to receive a good or bad rating while they are still in contact with the customer.

My main hypothesis is that the product and how the order was fulfilled influences customer review rating.

In this project, machine learning techniques will be applied to the dataset to predict customer review ratings.

Dataset:

Brazilian E-Commerce Public Dataset by Olist

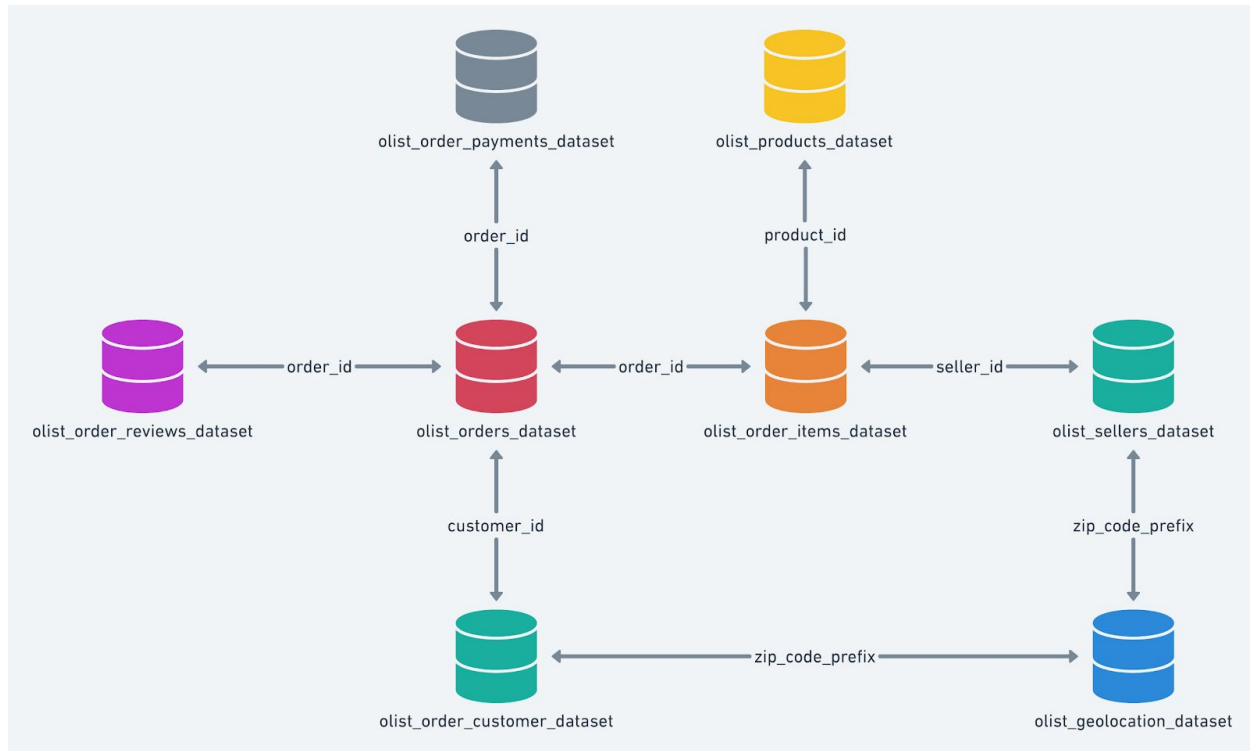
<https://www.kaggle.com/olistbr/brazilian-ecommerce>

This is a Brazilian ecommerce public dataset of orders made at Olist Store. The dataset has information of 100k orders from 2016 to 2018 made at multiple marketplaces in Brazil. It's features allows viewing an order from multiple dimensions: from order status, price, payment and freight performance to customer location, product attributes and finally reviews written by customers.

This is real commercial data, it has been anonymised, and references to the companies and partners in the review text have been replaced with the names of Game of Thrones great houses.

Data Schema:

The data is divided in multiple datasets for better understanding and organization.



Data descriptions:

Customers:

Customer_id : key to the orders dataset. Each order has a unique **customer_id**.

Customer_unique_id: unique identifier of a customer.

Customer_zip_code_prefix: first five digits of customer zip code

Customer_city: customer city name

Customer_state: customer state

Order Items

Order_id: order unique identifier

Freight_value: item freight value item (if an order has more than one item the freight value is splitted between items)

Order Reviews

review_comment_title: Comment title from the review left by the customer, in Portuguese.

Review_comment_message: Comment message from the review left by the customer, in Portuguese.

Review_creation_date : Shows the date in which the satisfaction survey was sent to the customer.

Review_answer_timestamp: Shows satisfaction survey answer timestamp.

Order

Order_approved_at :Shows the payment approval timestamp.

Order_delivered_carrier_date: Shows the order posting timestamp. When it was handled to the logistic partner.

Order_delivered_customer_date : Shows the actual order delivery date to the customer.

Order_estimated_delivery_date: Shows the estimated delivery date that was informed to customer at the purchase moment.

Products

product_photos_qty:number of product published photos

Product_weight_g: product weight measured in grams.

Product_length_cm: product length measured in centimeters.

Product_height_cm:product height measured in centimeters.

Product_width_cm: product width measured in centimeters.

Category Name Translation

Product_category_name: category name in Portuguese

Product_category_name_english : category name in English

Data Wrangling

This section describes data cleaning and wrangling methods applied to olist data set.

Firstly, the data which was in csv format was loaded and saved as a pandas dataframe.

Datetime columns:

order_purchase_timestamp

order_approved_at

order_estimated_delivery_date

order_delivered_customer_date

review_score

review_creation_date

review_answer_timestamp

All above columns with datetime values were converted from object to datetime datatype.

Product names:

The product names in the dataset were in portuguese. A new dataset `products_in_english` was loaded and mapped with current dataset to replace portuguese names to their English Meaning.

Order delivery columns:

No. of orders delivered after estimated delivery date was calculated by calculating the difference between `order_estimated_delivery_date` and `order_delivered_customer_date'`


If the value is negative, it means the order was delivered after estimated delivery date.
If positive, then the order was delivered before estimated delivery date.

The updated data frame was saved for further analysis.

Exploratory data analysis and visualization:

Once data wrangling and cleaning part was done, I started to explore the data to find correlation between different variables (features) from the dataset.

I tried initially plotting **Most bought product categories**



Bed Bath table,health beauty,sports are some categories that are bought most often by the customers.The difference in order count between bed bath table and health beauty is high.

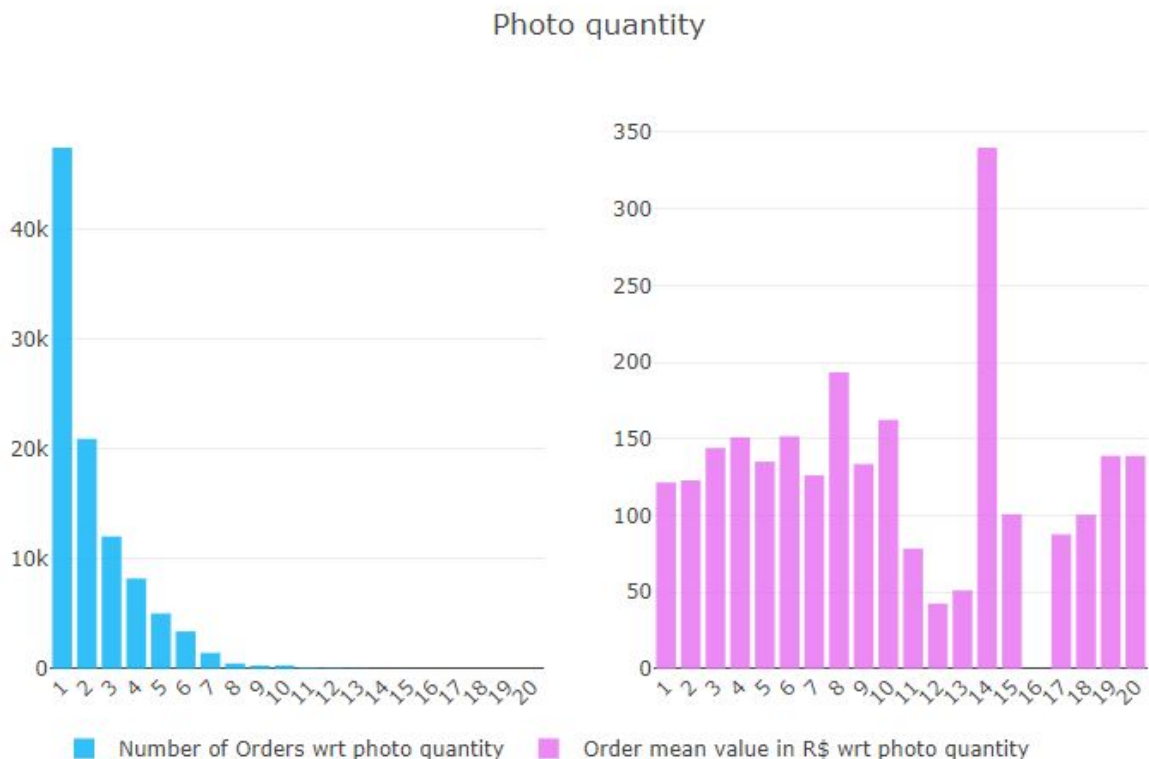
Number of products people usually order



Most of them have ordered only 1 product.The number of people ordering more than 2 items is very less.

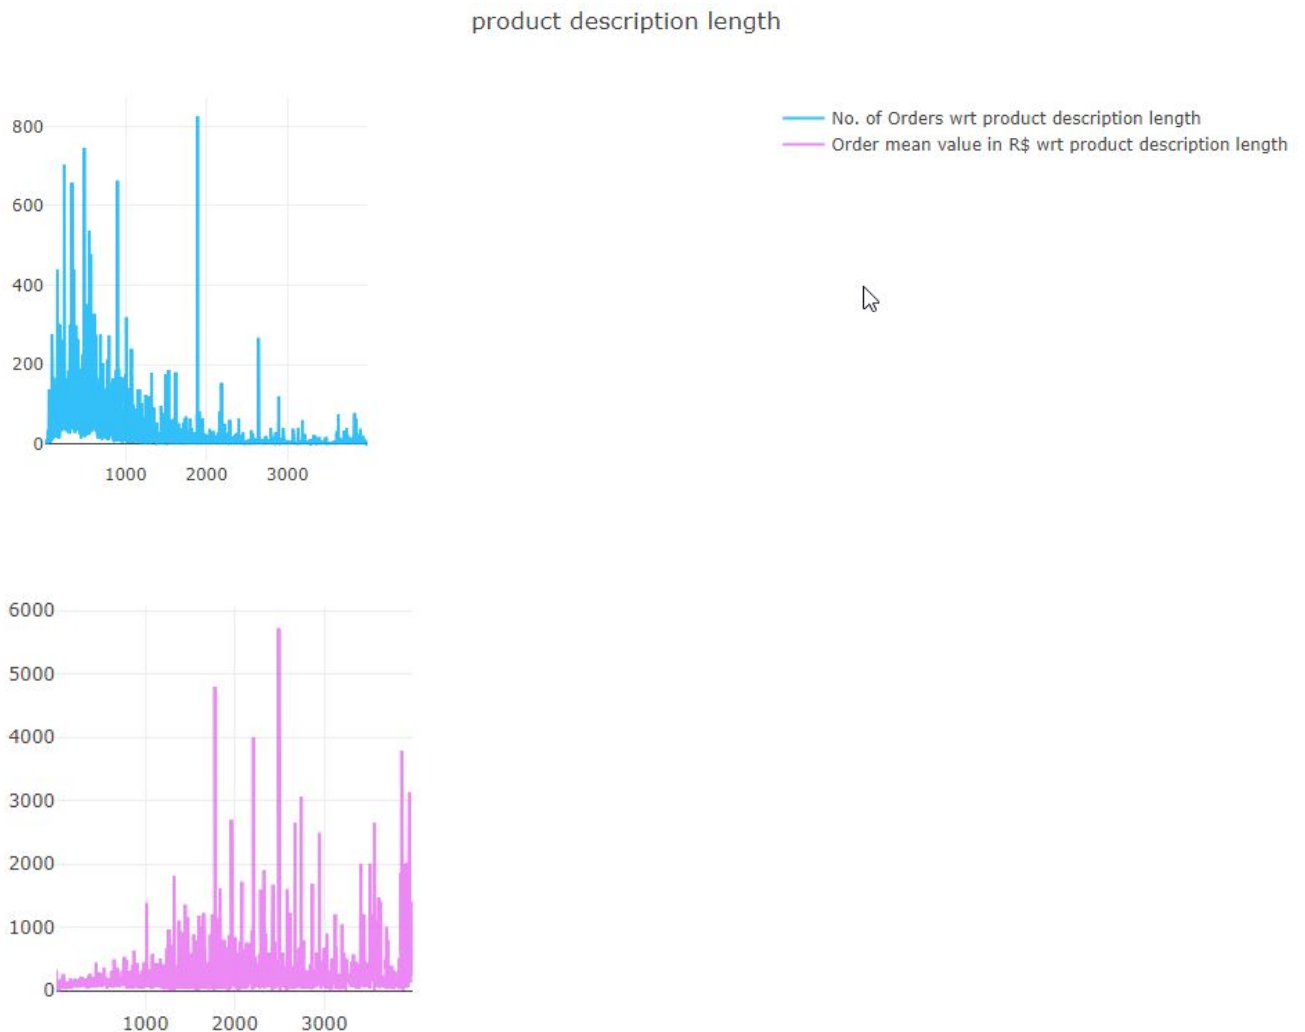
Number of photos and orders

Later, I wanted to check how the number of photos available on site are related to the orders.




There are more orders for the products with less than 2 photo quantities.

Product description length and order count:




There is no significant relation seen between the product description length and order count. Products with description length less than 2000 words are ordered frequently than the products with brief descriptions.

Product and freight values by date



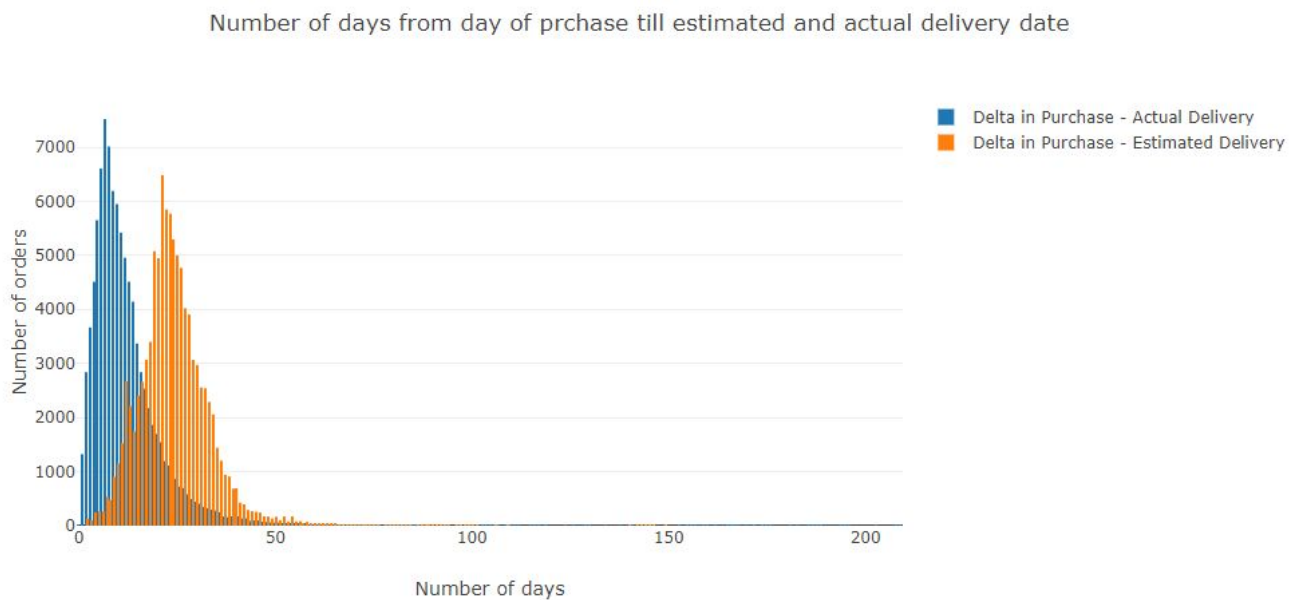
The freight value seems increasing with product value. There seems to be a similar trend in freight value and product value.

Money spent on each order/ Transaction value per order




There are more orders for less valued products so it seems that most often customer buy cheap product on Olist.

Let's check the average number of days between order and delivery



Most of the products seemed to be delivered before the estimated delivery date.


Explore review data



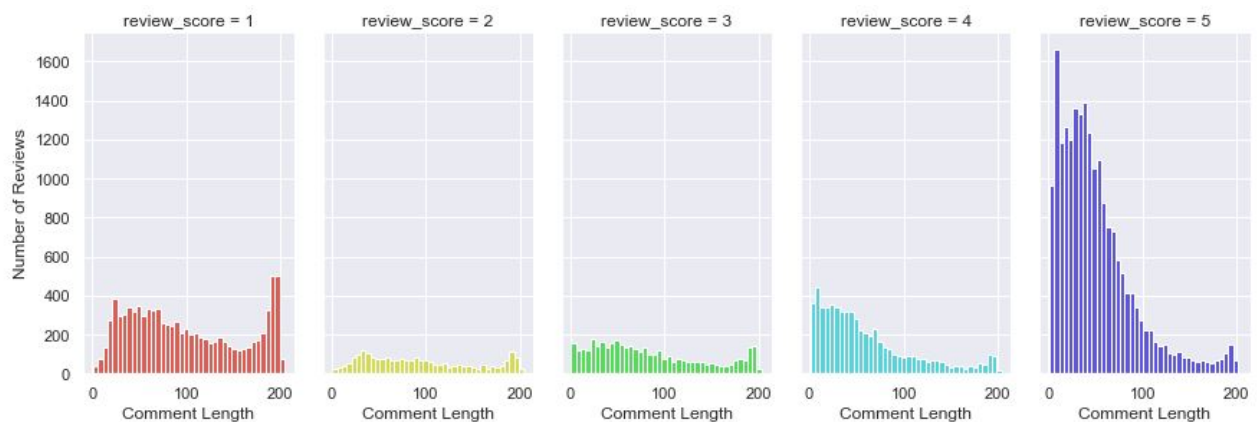
Here we can see the review score distribution. It is interesting to observe that there's more 1 star reviews than 2/3 stars reviews.

Capstone Project 2: Milestone report 1

If we plot the review score distribution of orders that do not have a *'delivered'* status, we can see that most of them have a 1 star rating.



Just for my curiosity I plotted the review score and the length of the review. From the plot, it looks like unhappy customers are more likely to write comments and there comment are long.



From visualizing data, I could address the question of whether or not the delivery delays cause customer dissatisfaction.

Statistical Inference:

From data exploration, I did not get much insight on the relationships between dependent variable ('review score') and independent variable but I could come up with a few hypothesis mentioned below.

Feature Hypothesis :

What are the expected relationships of independent variables with review score?

1. order_items_qty ("+") - if consumer gets more than one item from the same seller, it should mean that he/she knows the quality of the good. Therefore, increase in item quantity should increase review score
2. product_description_lenght ("+") - buyer having more information about buying product should have positive relationship with review score
3. product_photos_qty ("+")
4. product_name_lenght ('unknown')- shorter names sometimes can be ambiguous or can be easy to understand,so relationship with review score is unknown at this point
5. delivery_accuracy ("+") - item coming on time or earlier that it was described should have positive relationship with review score
6. order_products_value ("-") - more expensive perhaps means better quality or higher expectation towards order fulfilment;
7. order_freight_value ("-") - more freight means more expectation towards delivery accuracy

To check above, I ran an OLS on these features.

```
result.summary()
```

Dep. Variable:	review_score	R-squared:	0.069
Model:	OLS	Adj. R-squared:	0.068
Method:	Least Squares	F-statistic:	1197.
Date:	Sat, 09 Feb 2019	Prob (F-statistic):	0.00
Time:	23:46:46	Log-Likelihood:	-1.6176e+05
No. Observations:	97595	AIC:	3.235e+05
Df Residuals:	97588	BIC:	3.236e+05
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	4.0635	0.024	172.695	0.000	4.017	4.110
order_items_qty	-0.2464	0.009	-27.231	0.000	-0.264	-0.229
product_name_lenght	-0.0019	0.000	-4.616	0.000	-0.003	-0.001
product_description_lenght	4.038e-05	6.29e-06	6.424	0.000	2.81e-05	5.27e-05
product_photos_qty	0.0106	0.002	4.531	0.000	0.006	0.015
order_delivery_before_estimated_date_in_days	0.0317	0.000	79.803	0.000	0.031	0.033
order_products_value	-2.767e-05	2.2e-05	-1.261	0.207	-7.07e-05	1.54e-05

Omnibus:	17863.435	Durbin-Watson:	1.998
Prob(Omnibus):	0.000	Jarque-Bera (JB):	29327.189
Skew:	-1.288	Prob(JB):	0.00
Kurtosis:	3.757	Cond. No.	6.07e+03

Let's interpret the table. Overall the model is significant.

Prob(F-statistics)= 0.000 This is the p-value associated with the F-statistic ($F(6,97588) = 1197$). It is used in testing the null hypothesis that all of the model coefficients are 0.

This tells us that there is a significant difference in the group means.

Coef. : These are the values for the regression equation for predicting the dependent variable from the independent variable. The regression equation is presented in many different ways, for example:

$$Y_{\text{predicted}} = b_0 + b_1x_1 + b_2x_2 + b_3x_3 + b_4x_4$$

The column of estimates provides the values for b_0 , b_1 , b_2 , b_3 and b_4 for this equation.

order_items_qty:-0.2464 So for every unit increase in order item quantity, a 0.2464 unit decrease in review_score is predicted, holding all other variables constant.

product_name_lenght :-0.0019 for every unit increase in product_name_lenght, a 0.0019 unit decrease in review_score is predicted, holding all other variables constant.

product_description_lenght :4.038e-05 for every unit increase in product_description_lenght, a 4.038e-05 unit increase in review_score is predicted, holding all other variables constant.

product_photos_qty : 0.0106 for every unit increase in product_photos_qty, a 0.0106 unit increase in review_score is predicted, holding all other variables constant.

order_delivery_before_estimated_date_in_days: 0.0317 for every unit increase in order delivery accuracy, a 0.0317 unit increase in review_score is predicted, holding all other variables constant.

order_products_value :-2.767e-05 for every unit increase in product_name_lenght, a 2.767e-05 unit decrease in review_score is predicted, holding all other variables constant.

Looking at the p-values : $P > |t|$ – This column shows the 2-tailed p-values used in testing the null hypothesis that the coefficient (parameter) is 0. Using an alpha of 0.05:

The coefficient for all the variables is significantly different from 0 because their p-value is 0.000 or smaller than 0.05. This means here's no difference between the means and conclude that a significant difference does exist.

All of our independent variables are statistically significant which is a great news. But It's clear that we have to use more informative features to model this problem.

Feature Engineering:

The first feature I want to engineer is **Estimated Delivery Time in working days**. Gets the days between order approval and estimated delivery date. A customer might be unsatisfied if he is told that the estimated time is big.

After researching on olist site, it was found out that the most common carrier used for delivery does not deliver on Sunday so we have to consider this while calculating the difference between time of purchase and time of estimated delivery.

The second feature I want is **Estimated Delivery Time in working days**. Gets the days between order approval and delivered customer date. A customer might be more satisfied if he gets the product faster.

The third feature is **Delivery Time Delta in working days**. The difference between the actual and estimated date. If negative- order was delivered early, if positive - order was delivered late. A customer might be more satisfied if the order arrives sooner than expected, or unhappy if he receives after the deadline.

The fourth feature is **delayed**. Variable indicating if the order was delivered after the estimated date.

The fifth feature is **Average Product Value**. Cheaper products might have lower quality, leaving customers unhappy.

The sixth feature is **Order Freight Ratio** If a customer pays more for freight, he might expect a better service.

The final feature I want to look at is **Purchase Day of Week**. Does it affect how happy are the customers? I do not have a lot of intuition behind this step but I have a feeling this feature may be useful.

After creating new features, it's time to build a model. This will be covered as part of Machine learning section.